

## **FINAL CAPSTONE PROJECT REPORT**

**Project Title:** Enterprise-Grade Agentic Text-to-SQL System

**Submitted By:** Anubhav Agrawal

**Date:** 16<sup>th</sup> Jan, 2026

**Trainer:** TamalKant Mukherjee

---

## **TABLE OF CONTENTS**

- 1. Executive Summary**
  - 2. Problem Definition**
  - 3. System Architecture & Design**
  - 4. Technology Stack**
  - 5. Project Progress & Implementation Details**
  - 6. Runbook / Deployment Guide**
  - 7. Key Learnings & Challenges**
  - 8. Future Enhancements**
  - 9. Appendices (Screenshots & Visuals)**
- 

### **1. EXECUTIVE SUMMARY**

This project delivers an "Agentic Text-to-SQL System" designed to democratize data access within an enterprise. By leveraging Large Language Models (LLMs) via Groq (Llama 3.3) and a stateful orchestration engine (LangGraph), the system allows non-technical users to query complex SQL databases using natural language.

Key differentiators include a "Zero-Trust" PII filtering mechanism that protects sensitive customer data, a self-correcting agentic loop that fixes SQL errors automatically, and a real-time visualization layer built with Streamlit.

---

### **2. PROBLEM DEFINITION**

**The Challenge:** In modern enterprises, valuable data is locked in relational databases (SQL). Business stakeholders (Sales Managers, Executives) typically lack the SQL skills required to extract insights. This creates a bottleneck where every data request must go

through a technical data analyst, causing delays in decision-making. Thus, this will help millions of people to use data in their daily life with Simple natural language.

### Existing Solutions & Limitations:

- **Dashboards:** Static and cannot answer ad-hoc questions (e.g., "Why did sales drop yesterday?").
- **Basic Chatbots:** often "hallucinate" (make up data) or fail to understand complex schema relationships or wrong inputs.


**The Solution:** An intelligent, secure, and interactive interface that translates natural English questions into optimized SQL queries, executes them safely, and visualizes the results, all while scrubbing sensitive PII.

---

## 3. SYSTEM ARCHITECTURE & DESIGN

### 3.1 High-Level Architecture Diagram

The system follows a "Human-in-the-Loop" Agentic architecture.



### 3.2 Component Breakdown

1. **Orchestrator (LangGraph):** Manages the state machine. It decides if the agent should move forward to execution or loop back to fix an error based on the `should_retry` conditional edge.
  2. **Safety Layer (PII Filtering):** Implements a "Placeholder Swap" technique. Real emails (e.g., [alice@example.com](mailto:alice@example.com)) are swapped with tokens (`__EMAIL_1__`) before reaching the LLM, and swapped back only during execution.
  3. **Visualizer (Streamlit):** Automatically detects if the result is tabular or numeric and renders interactive Bar Charts and Dataframes.
- 

## 4. TECHNOLOGY STACK

- **Language:** Python 3.10+
  - **Frontend:** Streamlit
  - **LLM Provider:** Groq (Model: llama-3.3-70b-versatile)
  - **Orchestration:** LangGraph (StateGraph, Nodes, Edges)
  - **Database:** SQLite (Local relational DB with multiple tables)
  - **Data Handling:** Pandas (Data manipulation), Faker (Synthetic data generation)
  - **Visualization:** Graphviz (Schema diagrams), Streamlit Charts
- 

## 5. PROJECT PROGRESS REPORTS (Summary)

- **Phase 1: Database Design:** Created a robust schema with 4 tables (sales, customers, products, regions) and populated it with 200+ rows of realistic dummy data using Faker.
- **Phase 2: Baseline Agent:** Implemented a basic "Prompt-to-SQL" chain using LangChain and Groq.
- **Phase 3: Agentic Loops:** Migrated to LangGraph to handle errors. Added the `check_sql_node` for syntax verification and the `should_retry` loop for runtime error handling.
- **Phase 4: Advanced Security:** Implemented the "Zero-Trust" PII filter that redacts emails based on database presence.
- **Phase 5: UI & Visualization:** Built the Streamlit interface with the "Database Explorer" sidebar and dynamic charting capabilities.

---

## 6. RUNBOOK / DEPLOYMENT GUIDE

### Prerequisites

- Python 3.10 or higher installed.
- A valid Groq API Key.

### Step-by-Step Setup

1. Clone/Unzip Project: Extract the submitted final\_project.zip folder.

2. Install Dependencies: Open a terminal in the project folder and run:

```
pip install -r requirements.txt
```

3. Configure Environment: Open the backend.py file and ensure the API Key is set:

```
GROQ_API_KEY = (Used in the project)
```

4. Initialize Database: Generate the fresh dataset:

```
python database.py
```

*Output:* Database updated with Customers successfully.

5. Launch Application: Run the frontend server:

```
Python -m streamlit run main.py
```

6. Access UI: Open your browser to <http://localhost:8501>.

---

## 7. KEY LEARNINGS AND INSIGHTS

### Challenges Faced

1. **LLM Hallucinations:** The model initially tried to query columns like "dance\_date" which didn't exist.
  - **Solution:** We implemented "Schema Injection" in the system prompt and added a strict rule: "If column missing, return IMPOSSIBLE".
2. **Chatty LLM Responses:** The validation step failed because the LLM would write essays instead of just code.
  - **Solution:** Used Regex parsing to extract content solely between ```sql blocks.

3. **Data Privacy:** How to query specific customers without leaking their data to the cloud?
  - o **Solution:** The "Placeholder Swap" technique (sending \_\_EMAIL\_1\_\_ to the cloud, executing with real data locally).

## Takeaways

- **Agentic > Linear:** A self-correcting loop is infinitely more reliable than a single-shot prompt.
  - **Context should always be clear:** The quality of SQL generation is directly dependent on how clearly the Schema is presented in the prompt.
- 

## 8. DEMO SCREENSHOTS

(Please insert your screenshots in the spaces below)

Figure 1: The Chat Interface with Dynamic Visualization

[INSERT SCREENSHOT OF A QUERY LIKE "Top 3 Products" SHOWING THE BAR CHART]

Figure 2: The PII Security Alert Mechanism

[INSERT SCREENSHOT OF THE "Security Alert" MESSAGE WHEN ASKING FOR AN EMAIL]

Figure 3: The Self-Correction/Schema Diagram

[INSERT SCREENSHOT OF THE SIDEBAR SCHEMA DIAGRAM WITH ARROWS]

---

## 9. FINAL PRESENTATION DECK (Outline)

1. Title Slide: Enterprise Text-to-SQL Agent.
2. The Problem: "Data is hard to access."
3. The Solution: "Chat with your Database."
4. Live Demo:
  - o **Query 1 (Simple):** "Show me total sales."
  - o **Query 2 (Complex):** "Who are the top 3 customers in Europe?"
  - o **Query 3 (Safety):** "Is [admin@gmail.com](mailto:admin@gmail.com) in the database?" (Show Redaction).

- 5. Architecture: Show the LangGraph flow (Generate -> Check -> Execute).**
  - 6. Future Scope: Adding Voice Input and Multi-Database support.**
  - 7. Q&A.**
- 

**[End of Report]**